

SYDE 556/750

Simulating Neurobiological Systems

Lecture 8: Learning

Chris Eliasmith

October 23, 28, & 30

- ▶ Slide design: Andreas Stöckel
- ▶ Content: Terry Stewart, Andreas Stöckel, Chris Eliasmith



UNIVERSITY OF
WATERLOO

FACULTY OF
ENGINEERING



Learning


Definition

Learning is a directed change in synaptic weights \mathbf{W} while the network is active.

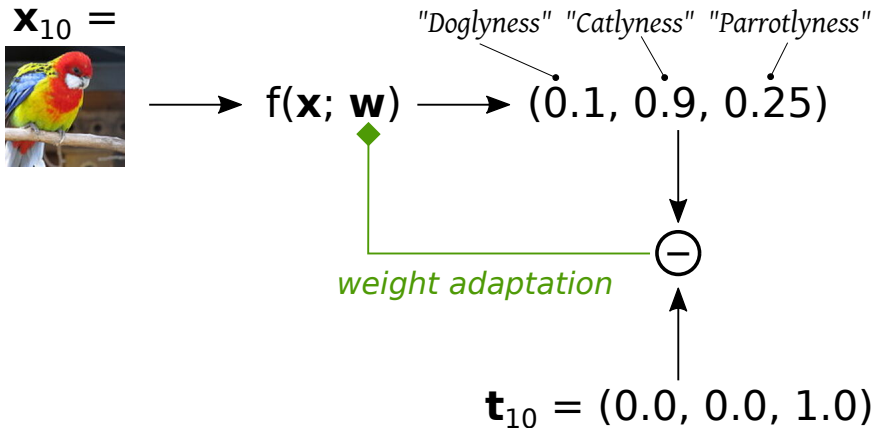
Learning is important because:

1. We might not know the function we want to compute at the beginning of a task.
2. The desired function might change over time.
3. The “optimal weights” we are solving for are not optimal.
4. We want to answer scientific questions about learning in nervous systems.


Supervised Learning



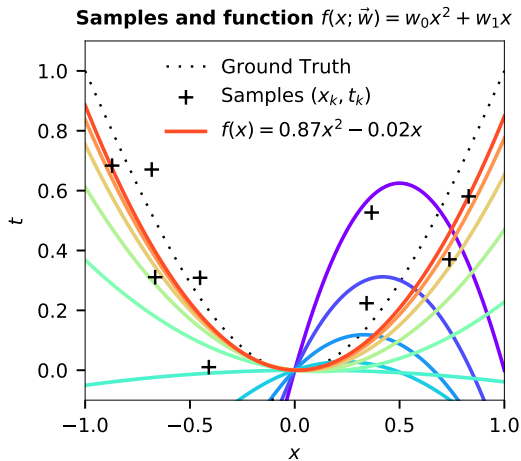
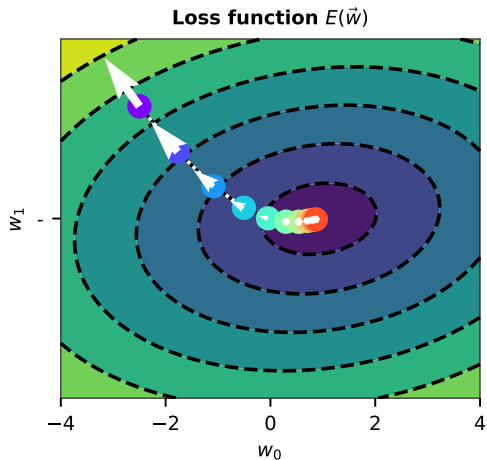
Supervised Learning – Training



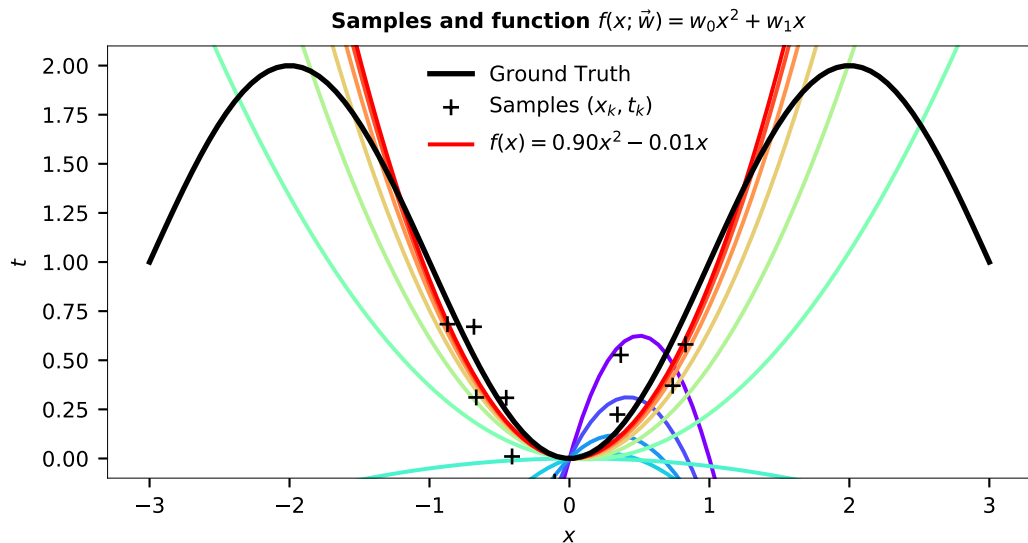
Supervised Learning – Inference



Gradient Descent – Example

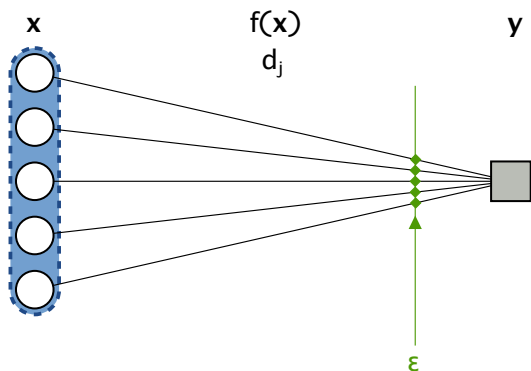


Supervised Learning – Generalisation



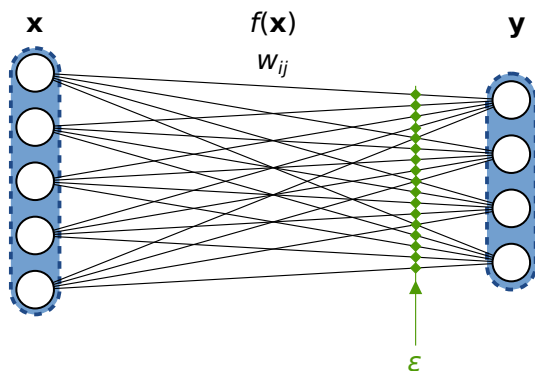
Learning Decoders and Learning Weights

Learning Decoders (Delta Rule)



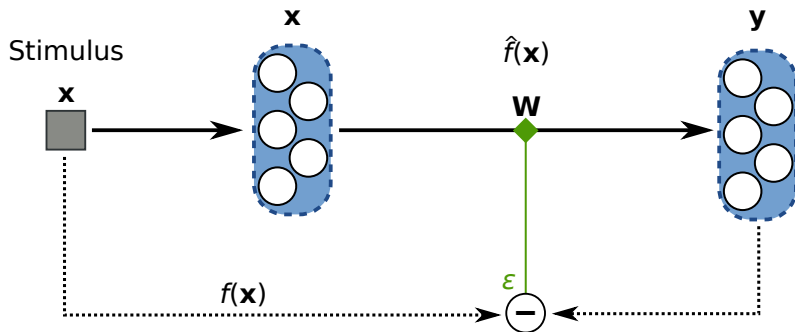
$$\Delta d_i = -\eta a_i(\mathbf{x}) \underbrace{(y(t) - y^d(t))}_{\varepsilon(t)}$$

Learning Weights (PES Rule)



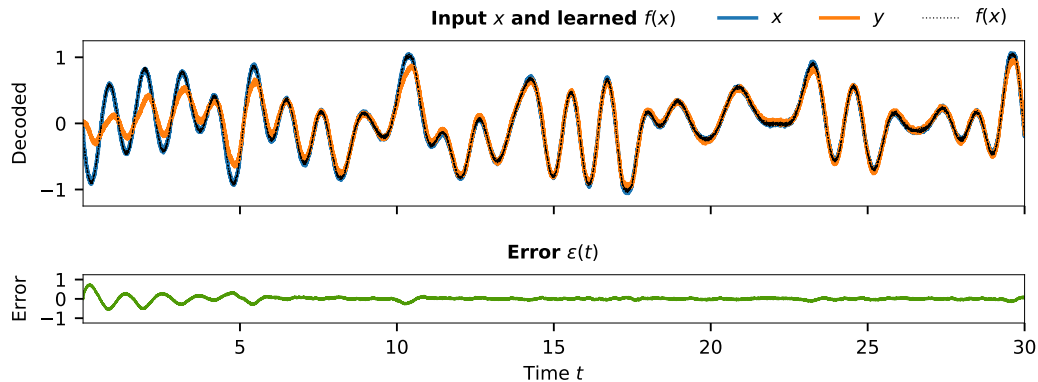
$$\Delta w_{ij} = -\eta a_i(\mathbf{x}) \left(\alpha_j \langle \mathbf{e}_j, \varepsilon(t) \rangle \right)$$

Example: Learning Functions (I)



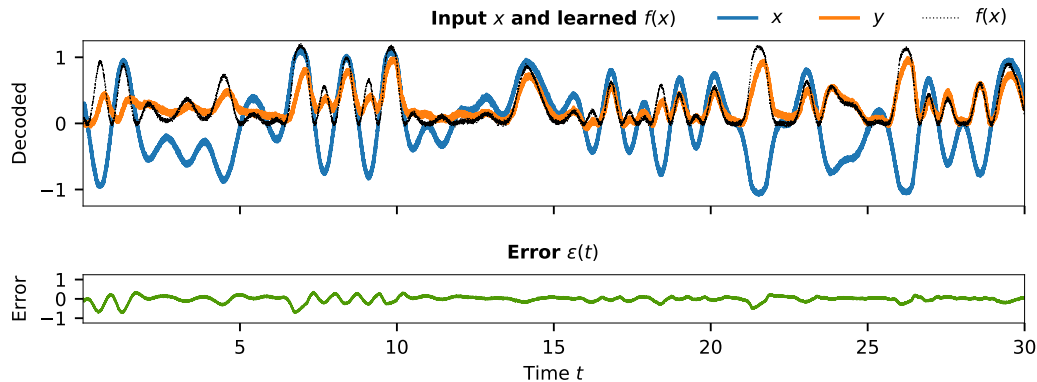
Example: Learning Functions (II)

Communication Channel $f(x) = x$



Example: Learning Functions (III)

Square $f(x) = x^2$




Works, but learns more slowly!


Where is the error signal $\varepsilon(t)$ coming from?

Example: Classical Conditioning (I)


Before conditioning:



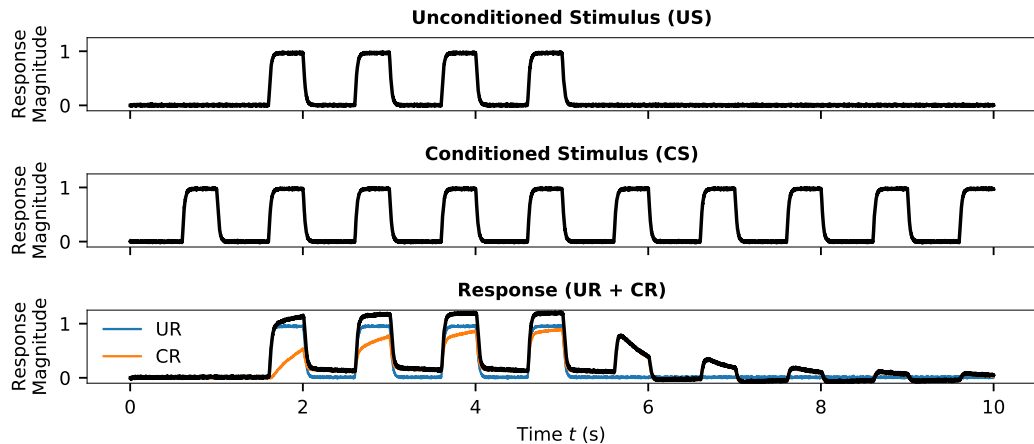
After conditioning:



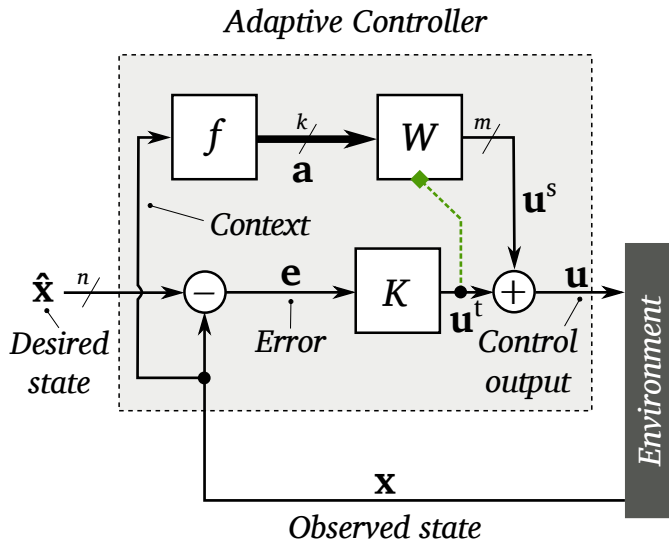
Example: Classical Conditioning (II)



Example: Classical Conditioning (III)



Example: Adaptive Controller



Unsupervised Learning

$\mathbf{x}_1 =$




$\mathbf{x}_2 =$




$\mathbf{x}_3 =$



$\mathbf{x}_4 =$



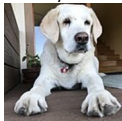
$\mathbf{x}_5 =$




$\mathbf{x}_6 =$



$\mathbf{x}_7 =$




$\mathbf{x}_8 =$



$\mathbf{x}_9 =$



$\mathbf{x}_{10} =$




$\mathbf{x}_{11} =$




$\mathbf{x}_{12} =$




$\mathbf{x}_{13} =$



$\mathbf{x}_{14} =$



$\mathbf{x}_{15} =$




$\mathbf{x}_{16} =$




$\mathbf{x}_{17} =$




$\mathbf{x}_{18} =$




$\mathbf{x}_{20} =$




$\mathbf{x}_{21} =$



$\mathbf{x}_{22} =$



$\mathbf{x}_{23} =$



$\mathbf{x}_{24} =$



Unsupervised Learning – Training




Unsupervised Learning – Inference




Dimensionality Reduction Example: Faces


Face Database

- ▶ 84 images of 12 women with 7 different expressions
- ▶ Normalised eye location
- ▶ 45×60 pixels (2700 dimensions)
- ▶ Greyscale



Autoencoder





PCA in Python



```
def PCA(X): # X: N x d matrix
    N, d = X.shape
    X_cen = X - np.mean(X, axis=0)
    C = (X_cen.T @ X_cen) / (N - 1)
    L, V = np.linalg.eigh(C) # "eigh" faster than "eig" for symmetric matrices
    return V.T[::-1, :] # d x d matrix
```




```
def PCA_SVD(X): # X: N x d matrix
    return np.linalg.svd(X - np.mean(X, axis=0))[2]
```

PCA Example: Source Images


Face Database

- ▶ 84 images of 12 women with 7 different expressions
- ▶ Normalised eye location
- ▶ 45×60 pixels (2700 dimensions)
- ▶ Greyscale




PCA Example: Eigenfaces



Identity Basis




Principal Components




PCA Example: Face Spaces






PCA Example: Sparse Vectors



PCA Example: Modifying the Latent Space



Limitations of PCA: Classifying Two Groups



Limitations of PCA: Metaphorical Illustration

Optimally extracted manifold X'



Dataset X



PCA X'



PCA

Projects ("squashes") data
onto a hyperplane

Hebbian Learning


When an axon of cell A is near enough to excite a cell B and repeatedly or persistently takes part in firing it, some growth process or metabolic change takes place in one or both cells such that A 's efficiency, as one of the cells firing B , is increased.

— Donald O. Hebb, “The Organization of Behaviour”, 1949

Hebbian Learning


$$\Delta w_{ij} = \eta a_i a_j$$

Example: Normalised Hebbian Learning



Learning an encoder \mathbf{e} , with $\|\mathbf{e}\| = 1$, 10000 steps, $\eta = 0.2 \times 10^{-4}$, $\Delta \mathbf{e} = \eta(\mathbf{x} \circ \mathbf{e})$

Spike-Time Dependent Plasticity



Conclusion

Supervised Learning

- ▶ Find \mathbf{w} such that $f(\mathbf{x}_k; \mathbf{w}) \approx \mathbf{t}_k$
- ▶ *Hope:* $f(\mathbf{x}_k; \mathbf{w}) \approx f_{\text{GT}}(\mathbf{x}_k)$
- ▶ Use gradient descent to find \mathbf{w}
- ▶ Delta, PES learning rules
- ▶ Modulatory synapses in the brain

Unsupervised Learning

- ▶ Dimensionality reduction $f(\mathbf{x}_k) = \lambda_k$
- ▶ *Hope:* latent dimensions λ are “meaningful”
- ▶ Autoencoders (nonlinear), PCA (linear)
- ▶ Hebbian learning \Rightarrow learns PCA

Image sources

Title slide

Page from “Liber ethicorum des Henricus de Alemannia”. Title: “Henricus de Alemannia con i suoi studenti” (Henricus of Germany with his students), second half of 14th century.
From Wikimedia.