

Spatial Semantic Pointers (SSPs)

Chris Eliasmith
SYDE 556/750



Spatial Semantic Pointers

- Semantic pointers represent standard discrete structures (lists, trees, etc.)
- SSPs allow recurrent convolutions to have fractional powers

$$B^k = \underbrace{B \circledast B \circledast \dots \circledast B}_{B \text{ appears } k \text{ times}}$$

- Compute fractional k in Fourier space

$$B^k = \mathcal{F}^{-1} \left\{ \mathcal{F} \{B\}^k \right\}, \quad k \in \mathbb{R}$$

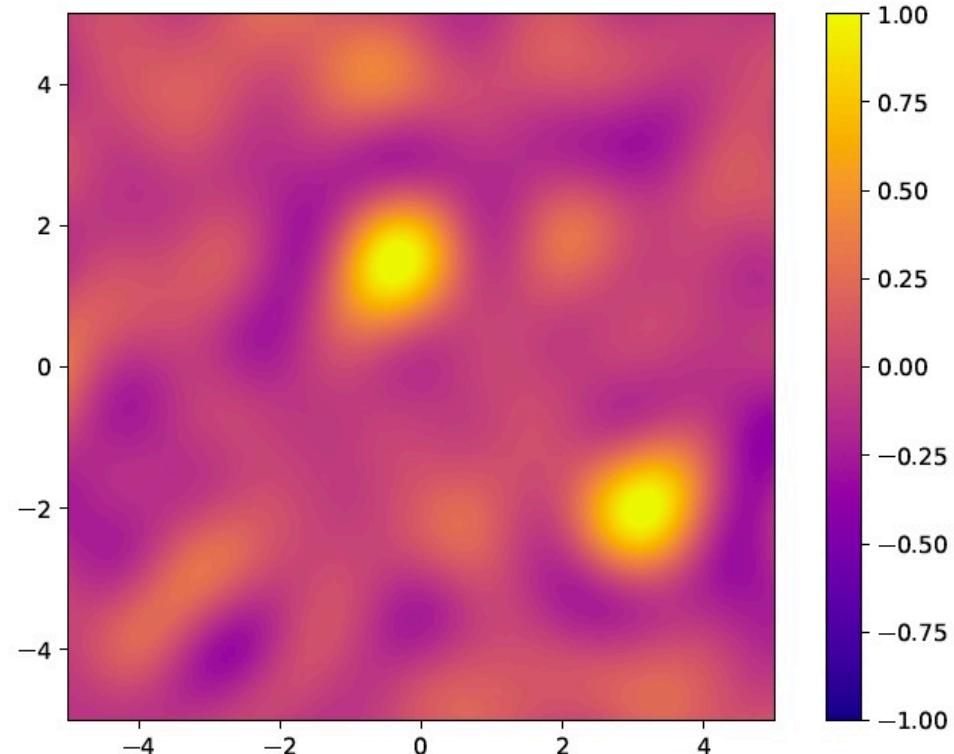
$$S(x, y) = X^x \circledast Y^y = \mathcal{F}^{-1} \{ \mathcal{F} \{X\}^x \odot \mathcal{F} \{Y\}^y \}$$

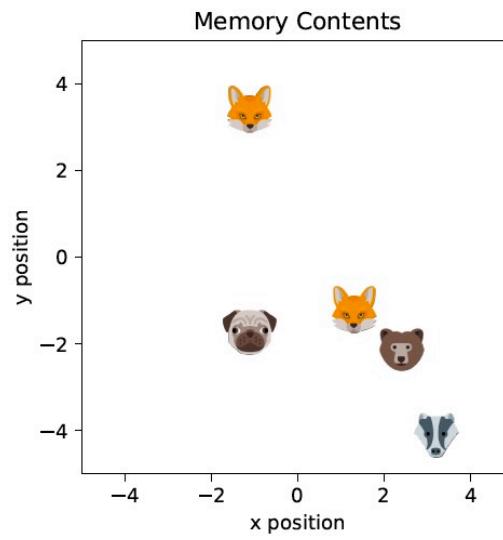
Spatial Semantic Pointers

- Represent continuous space (Clifford torus)

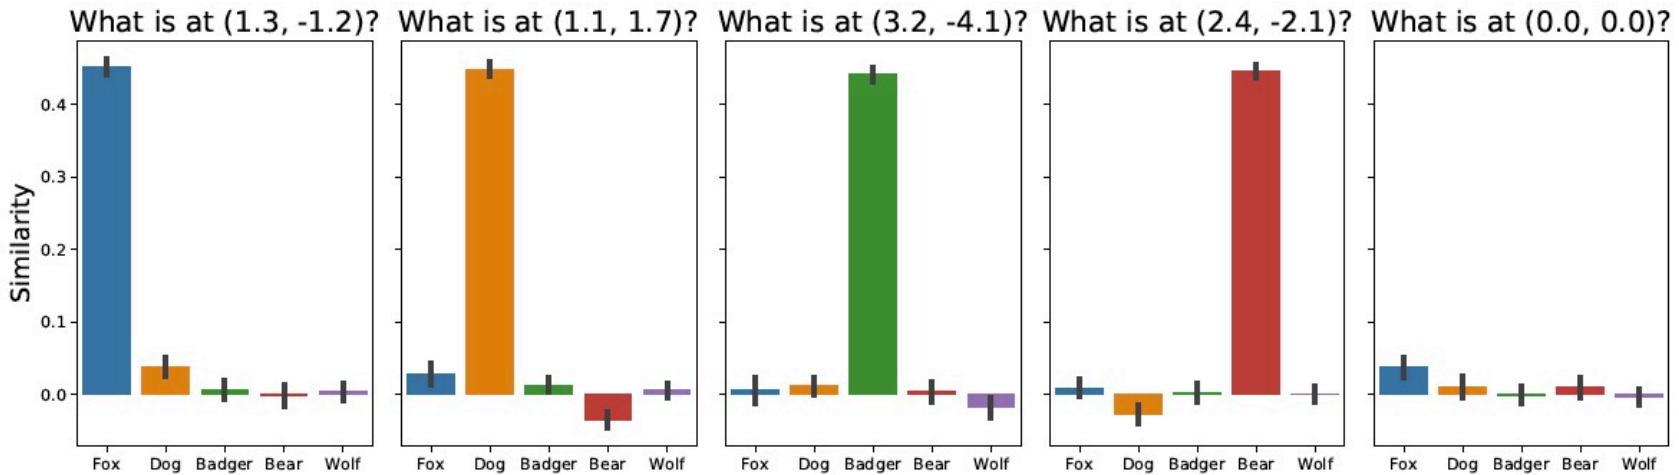
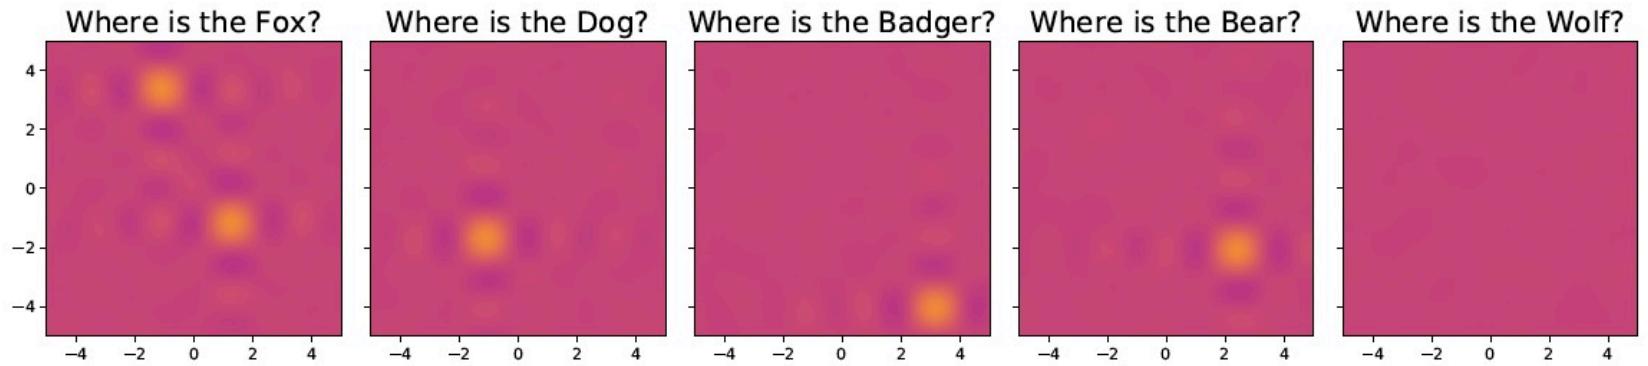
$$S(x, y) = X^x \circledast Y^y$$

- Heat map to visualize vector contents
- Dot product between SSP at every possible position and S





$$M = \sum_{i=1}^m OBJ_i \circledast S_i$$



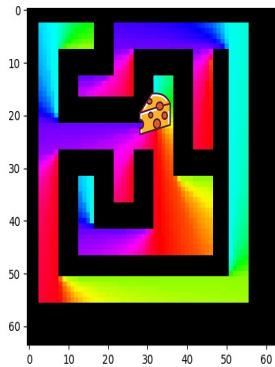
Spatial Manipulations

Desiderata	Accuracy	
	Non-Neural	Neural
Query single object	99.1%	95.7%
Query missing object	99.4%	96.7%
Query location	97.3%	94.7%
Query duplicate object	97.4%	95.3%
Query Region	90.4%	73.5%
Slide single object in group (all objects)	75.7%	67.3%
Slide single object in group (moved object)	100.0%	100.0%
Slide whole group	97.8%	96.7%
Readout x-y location from SSP	95.7%	94.1%
Construct SSP from x-y loca- tion	100.0%	99.0%

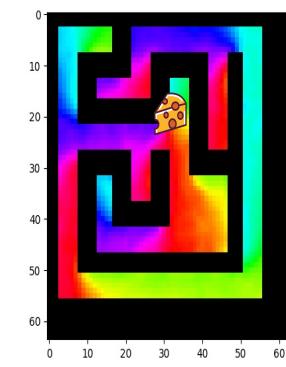
Navigating to Goal

Single layer MLP,
same for each
encoding (tried
many more).

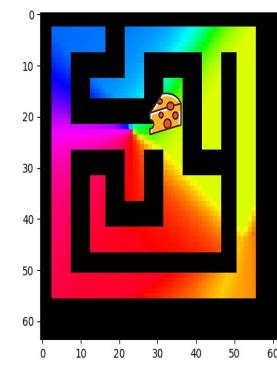
Ground Truth
RMSE: 0



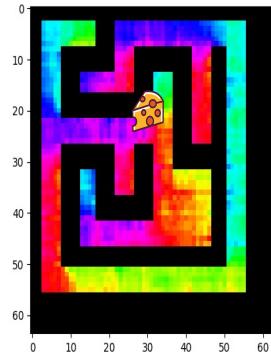
SSPs
RMSE: 0.0529



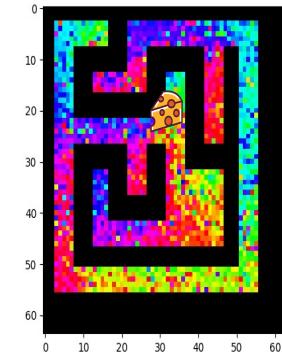
Random Projection
RMSE: 0.5351



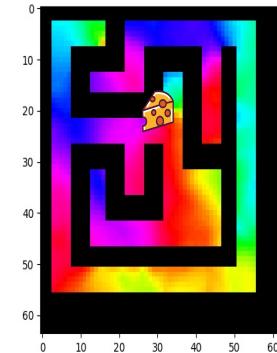
Trig Functions
RMSE: 0.1350



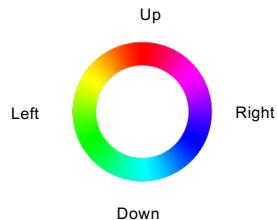
Random Mapping
RMSE: 0.2580



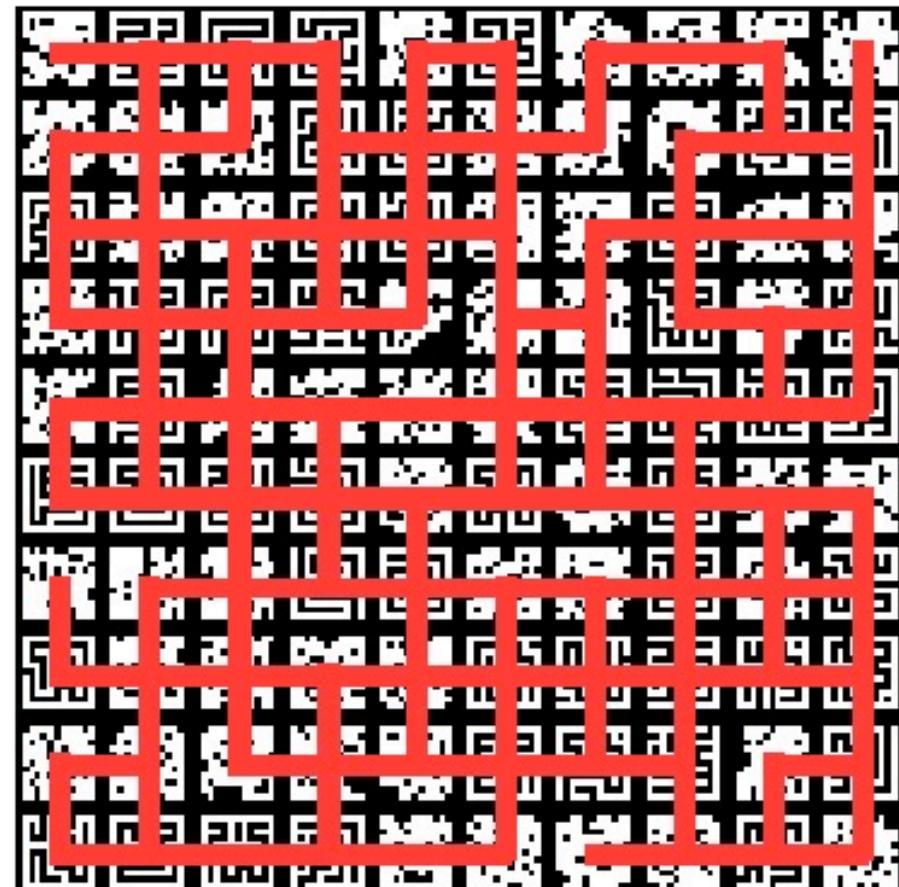
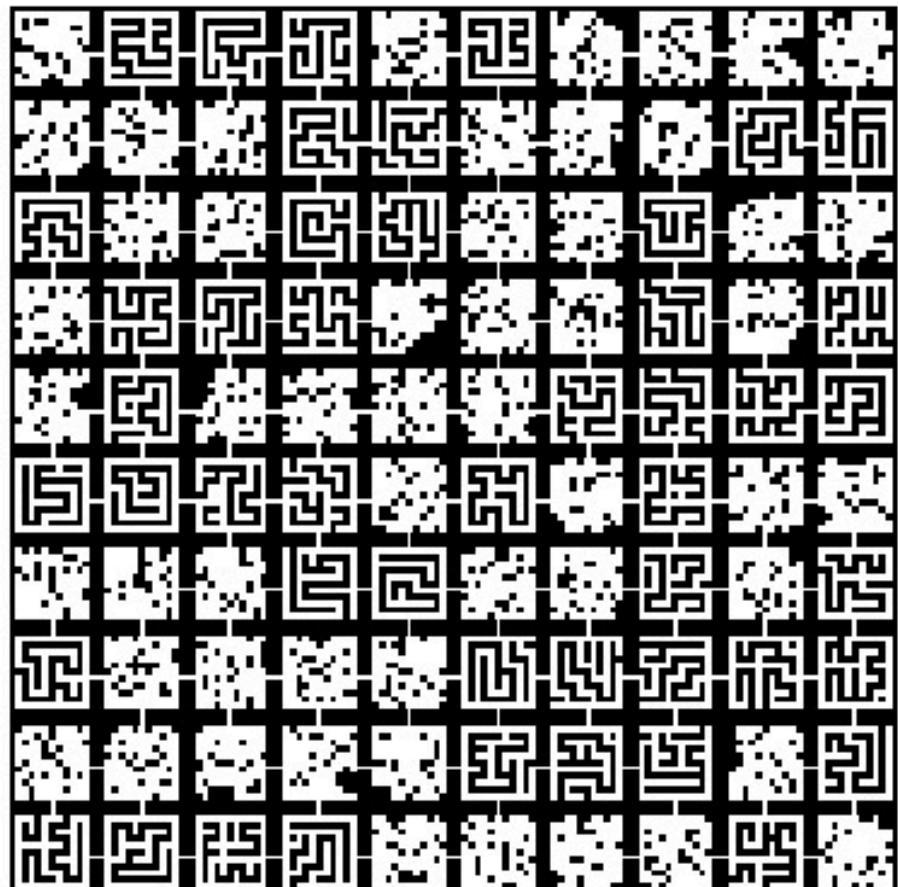
2D Coordinates
RMSE 0.1984



Legend

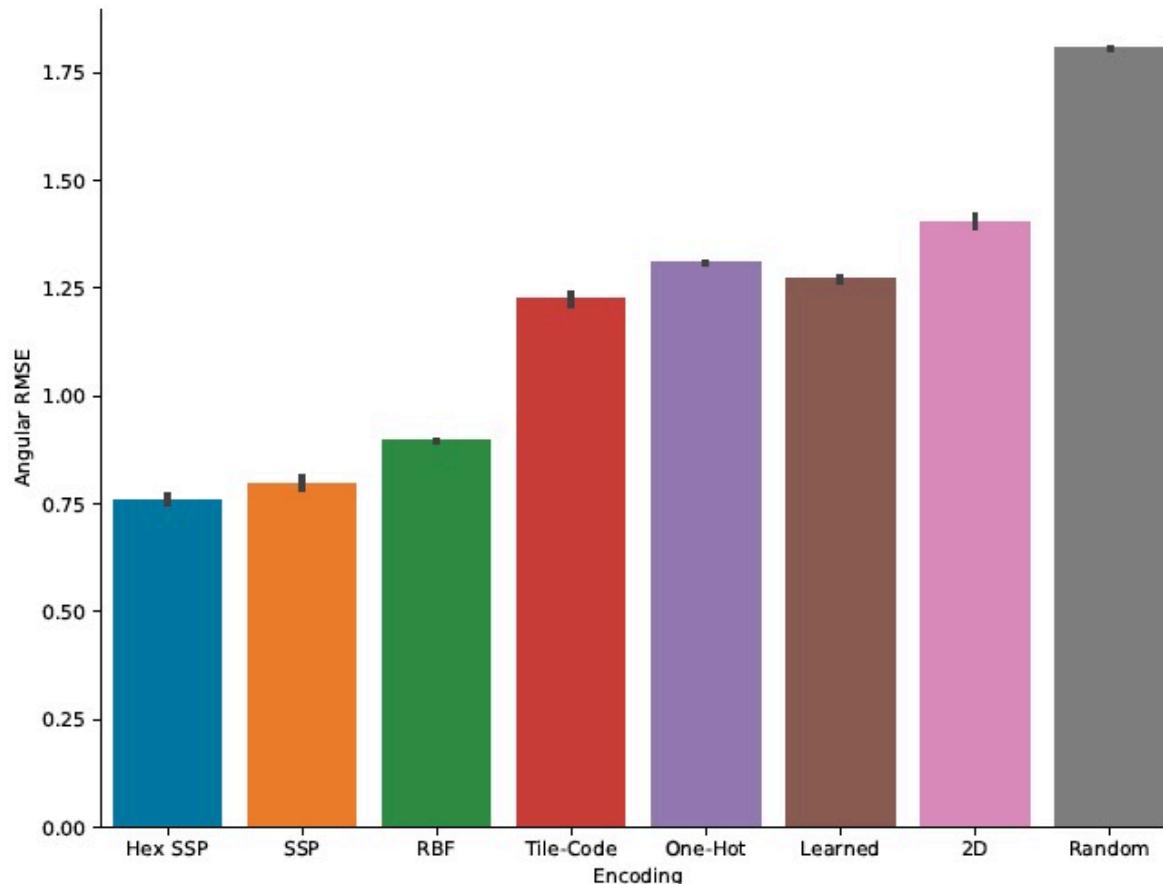


Scaling



Example 10x10 joined, hierarchical environment

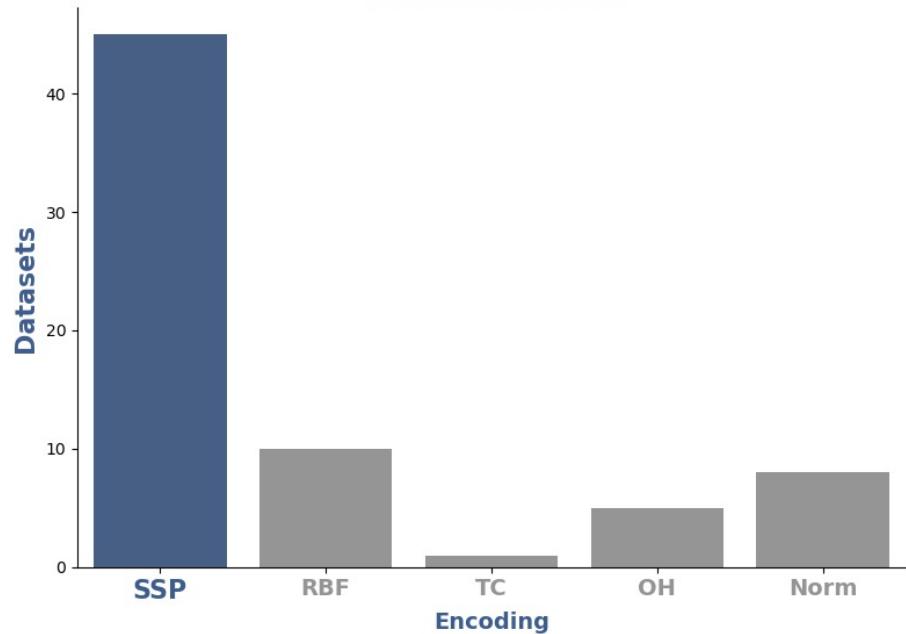
Scaling



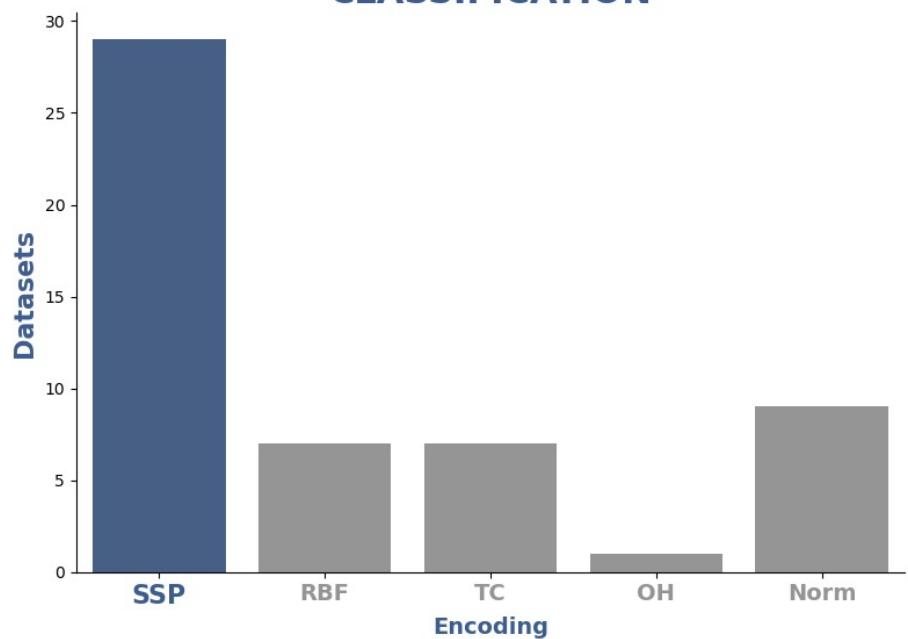
Performance of different encoders on 10x10
large maze from any point to any other

Good for General ML

REGRESSION



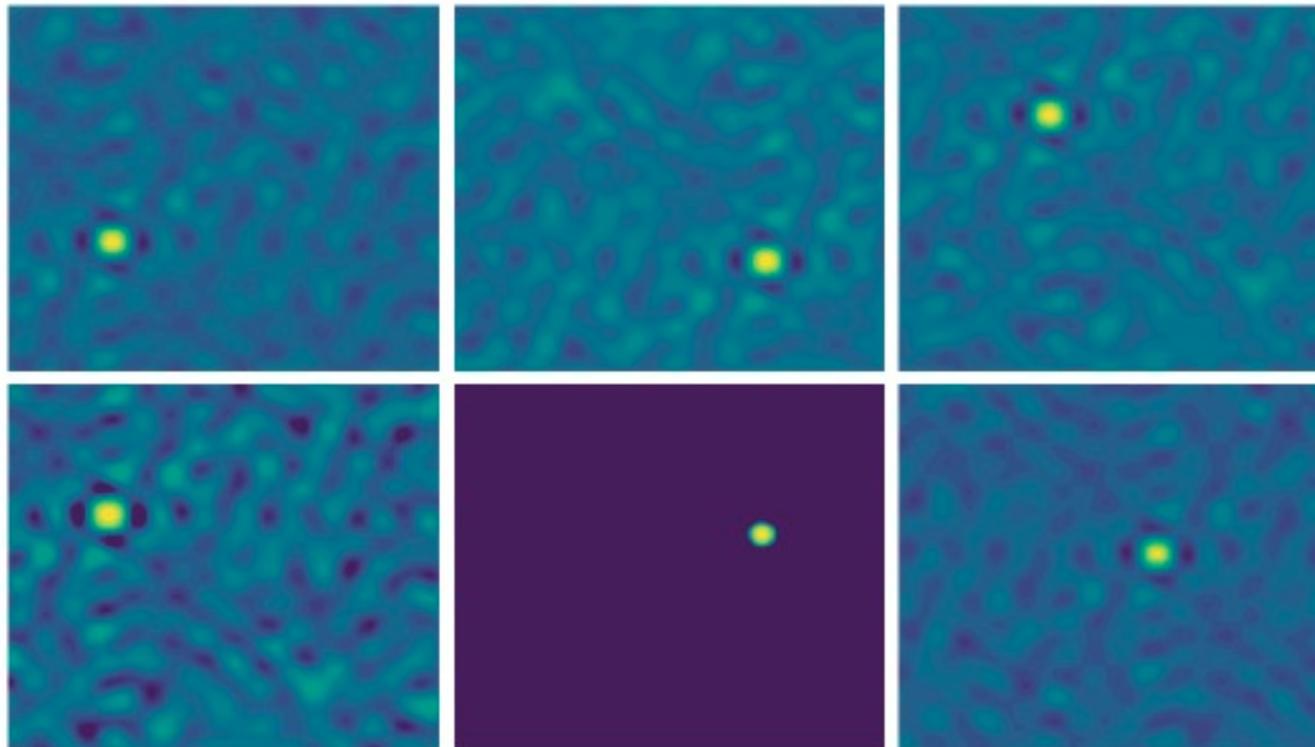
CLASSIFICATION



SSPs are more accurate on a large majority of 122 standard ML benchmarks.

How to Choose Axis Vectors

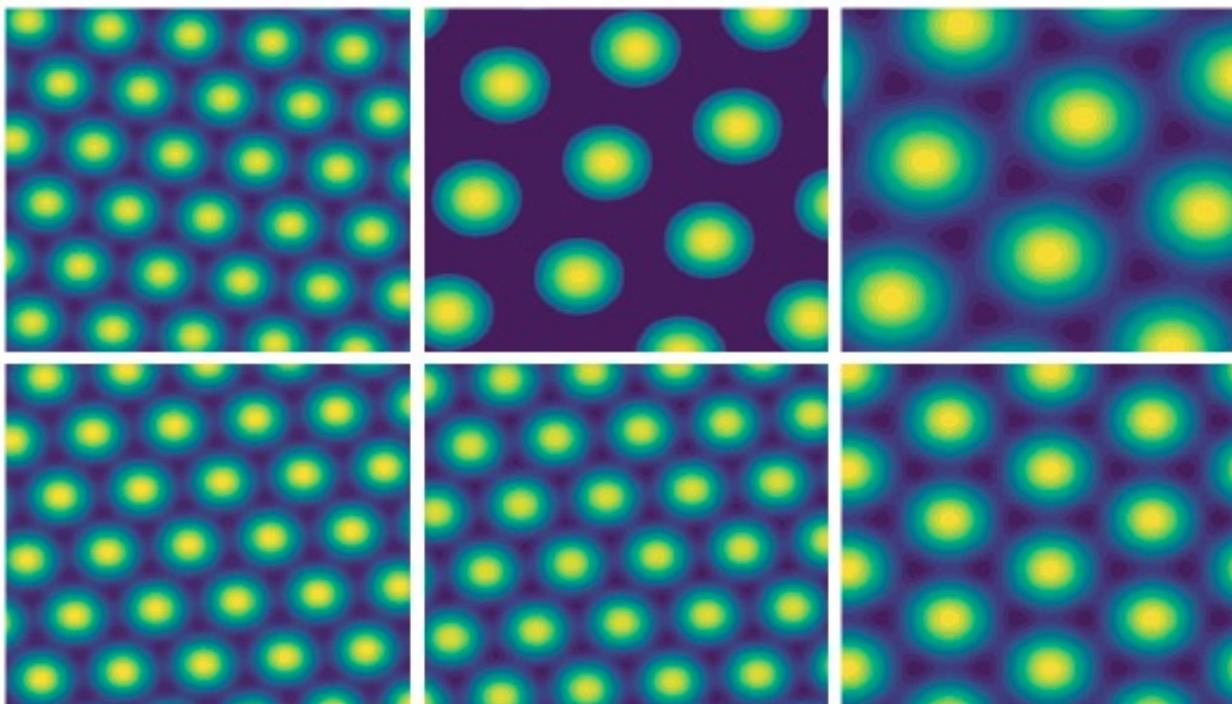
- Randomly (examples til now)



Tuning curves of neurons with random axis vectors and evenly tiled SSPs as encoders

How to Choose Axis Vectors

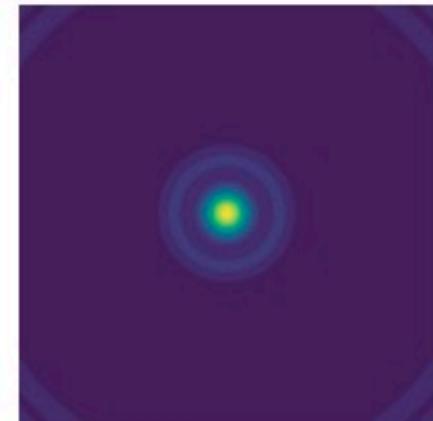
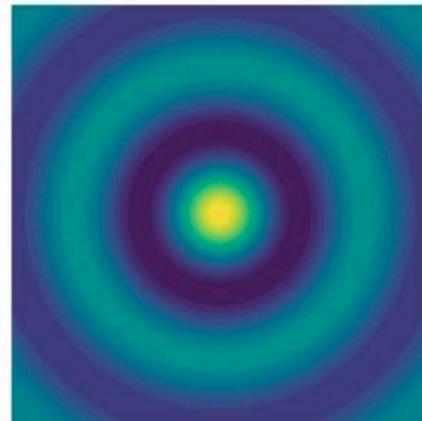
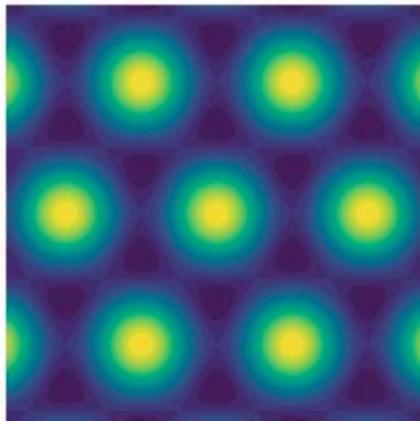
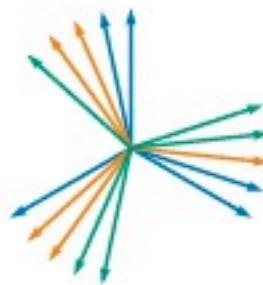
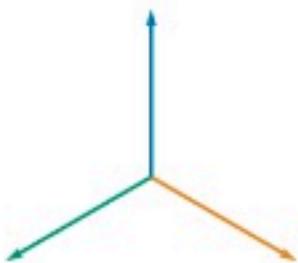
- With plane wave structure



Tuning curves of neurons with structured axis vectors and encoders picking out plane waves

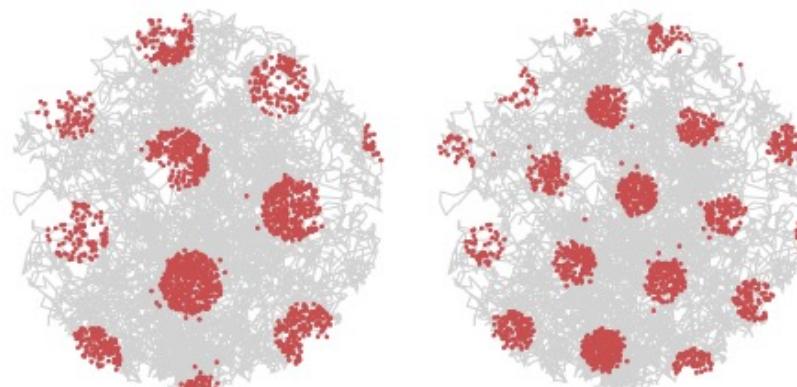
How to Choose Axis Vectors

- Sums of planar waves



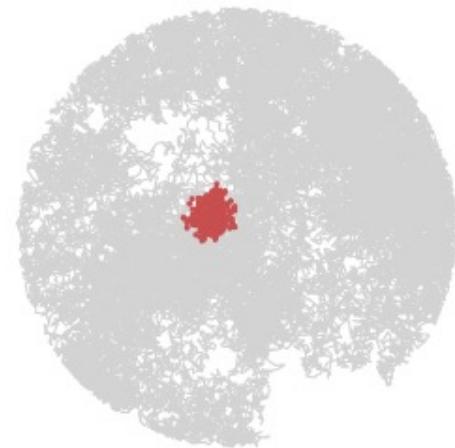
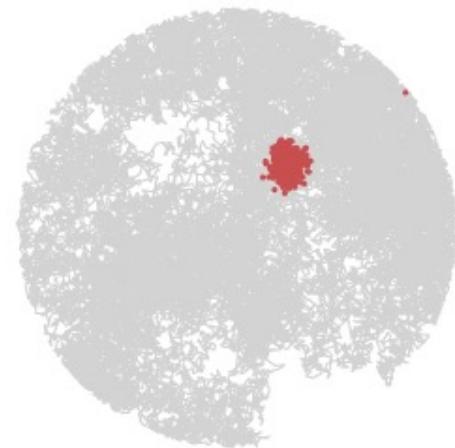
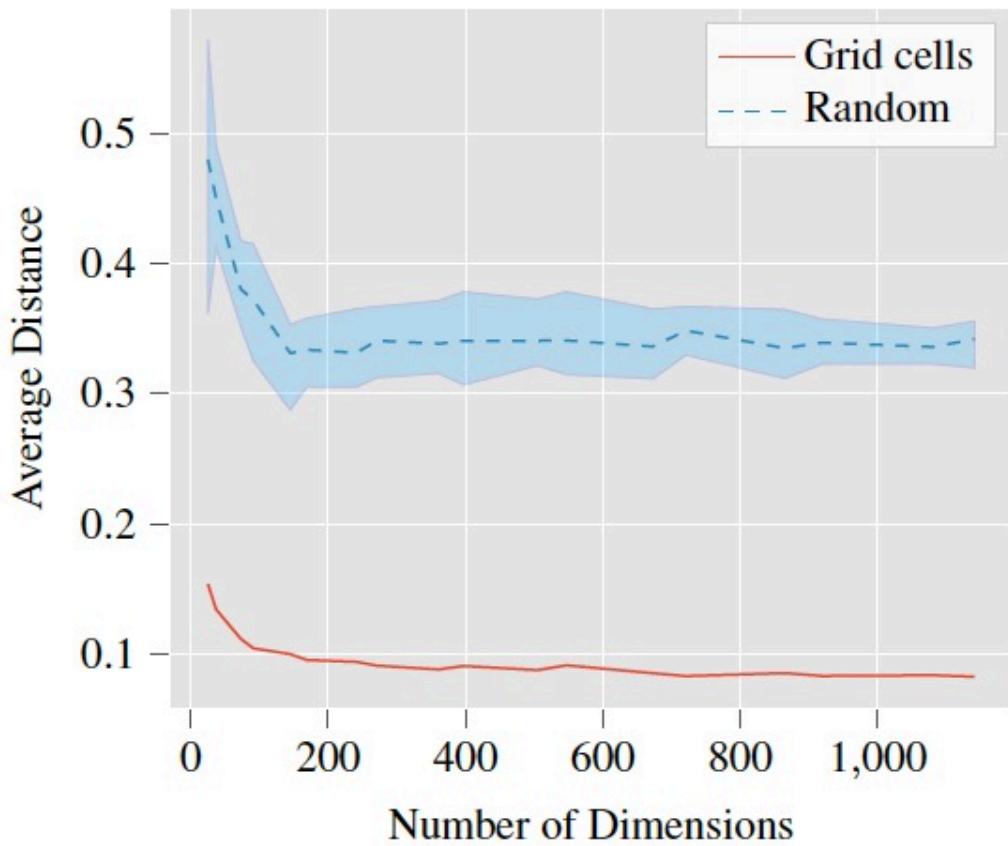
Grid Cells

- With plane wave structure, spiking neurons give grid cell responses
- We can combine them to get place cells (with standard NEF decoders)



SSP Grid cells

Place Cells



Mathematical properties

- Euclidean space is preserved on a high-d torus

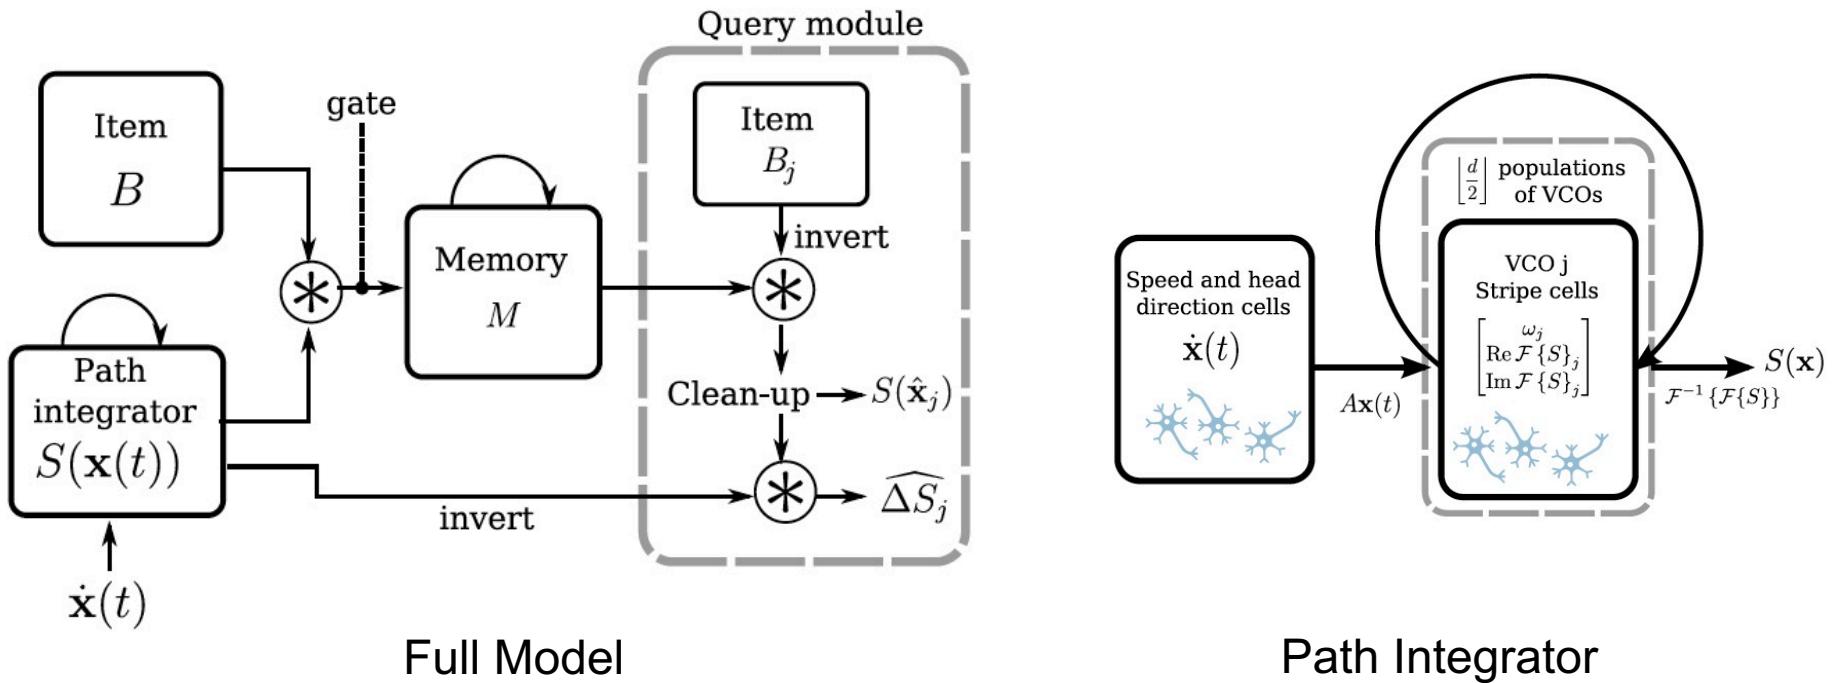
$$S(x_1, y_1) \circledast S(x_2, y_2) = S(x_1 + x_2, y_1 + y_2)$$

- E.g., $S(x, y) \circledast S(\Delta x, \Delta y) = X^{x+\Delta x} \circledast Y^{y+\Delta y}$
- We get the benefits of high-d representation, with accurate low-d Euclidean representation

Cognitive SLAM

- SLAM with complex features at certain spatial locations
- Start with no knowledge, bind vector descriptions to particular location in space
- Combines spatial and ‘symbol’ repn in neural network
- Semantic map as opposed to standard ‘image registration’ map

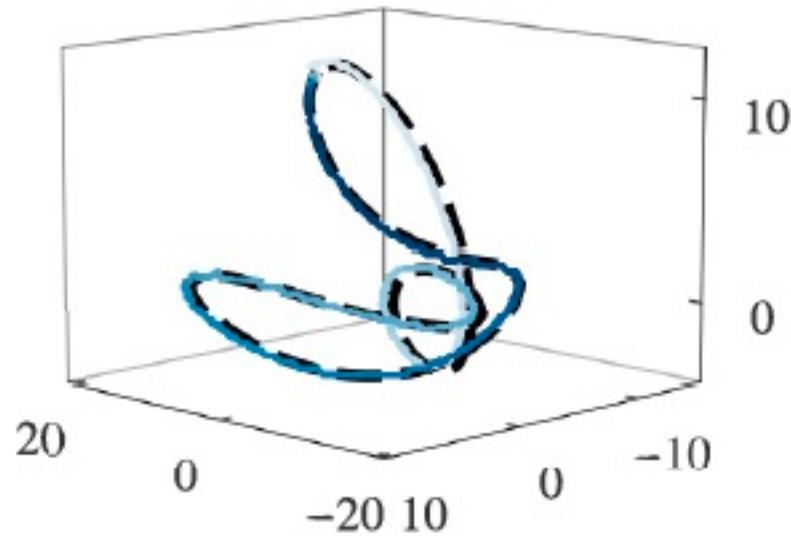
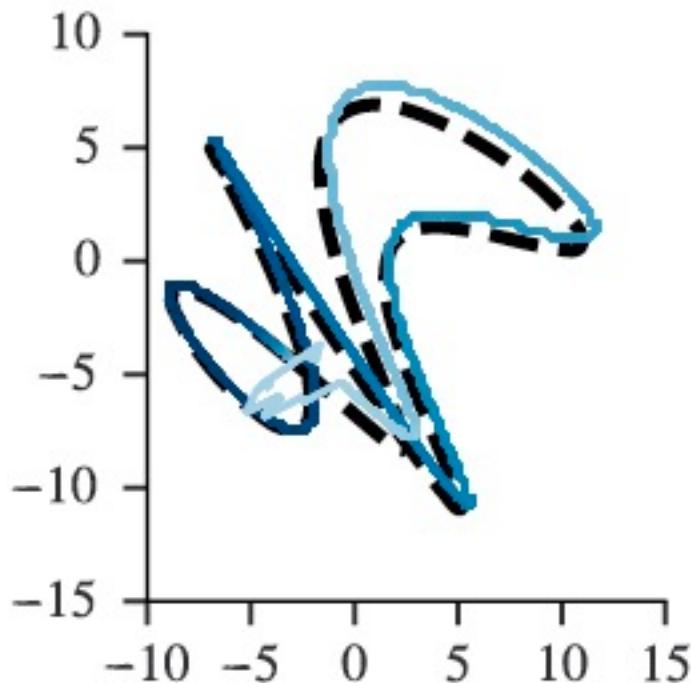
Cognitive SLAM Model



- Path integrator tracks ego position from velocity
- SLAM model learns env map, outputs allo- and ego-centric position

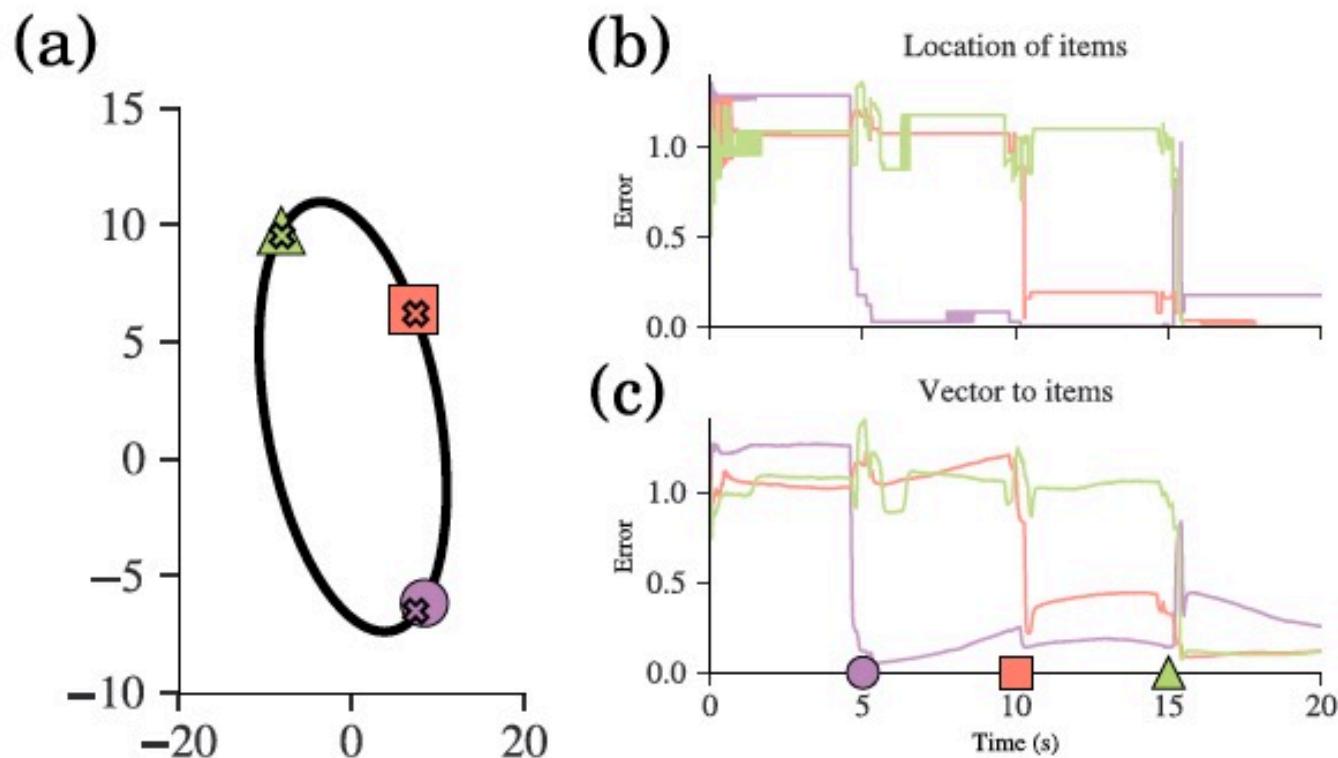
Path Integration

- One-minute long paths in 2D and 3D, spiking network



Cognitive SLAM

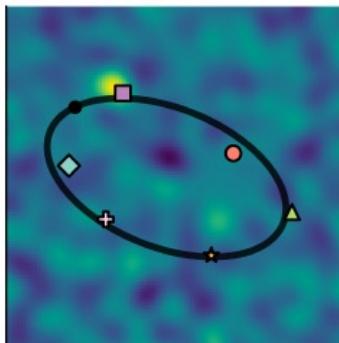
- Learns env by end of 20s path, full spiking
- Scaling up; symbols ‘over’ space



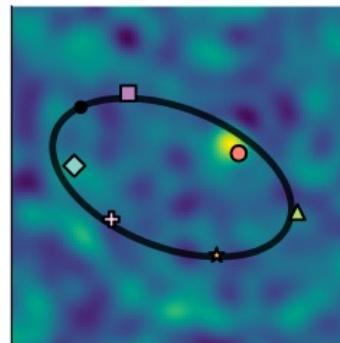
Cognitive SLAM

- Learns maps in LTM bound to position

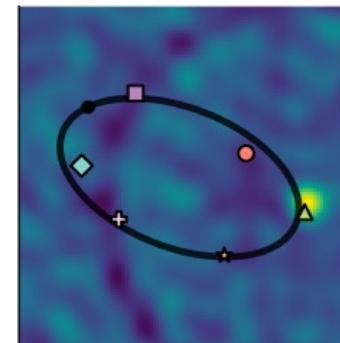
SQUARE



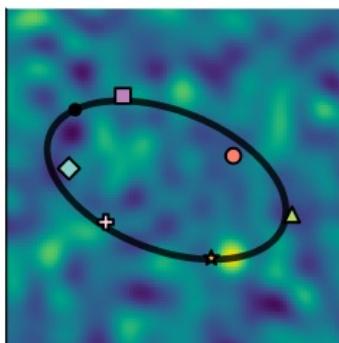
CIRCLE



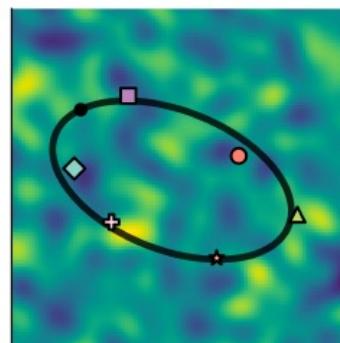
TRIANGLE



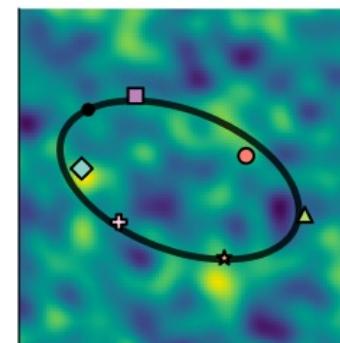
STAR



PLUS

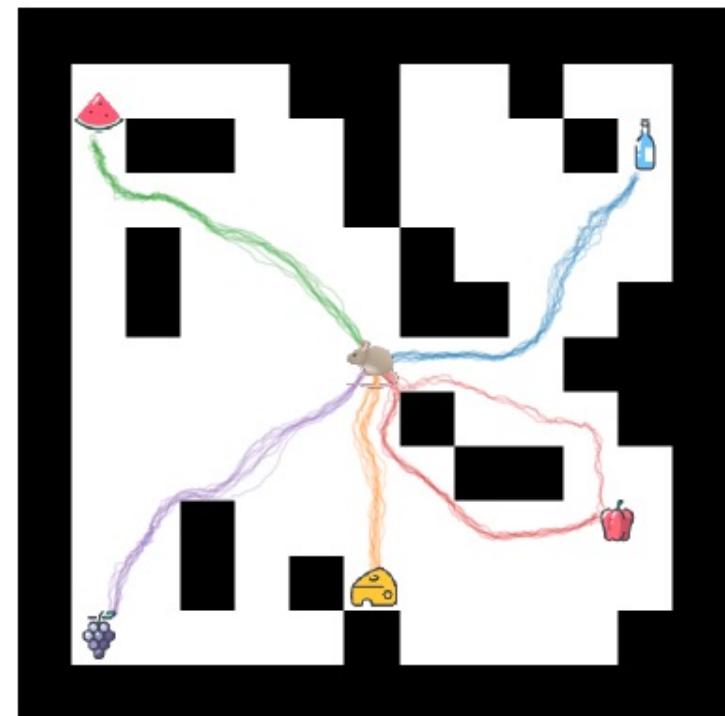
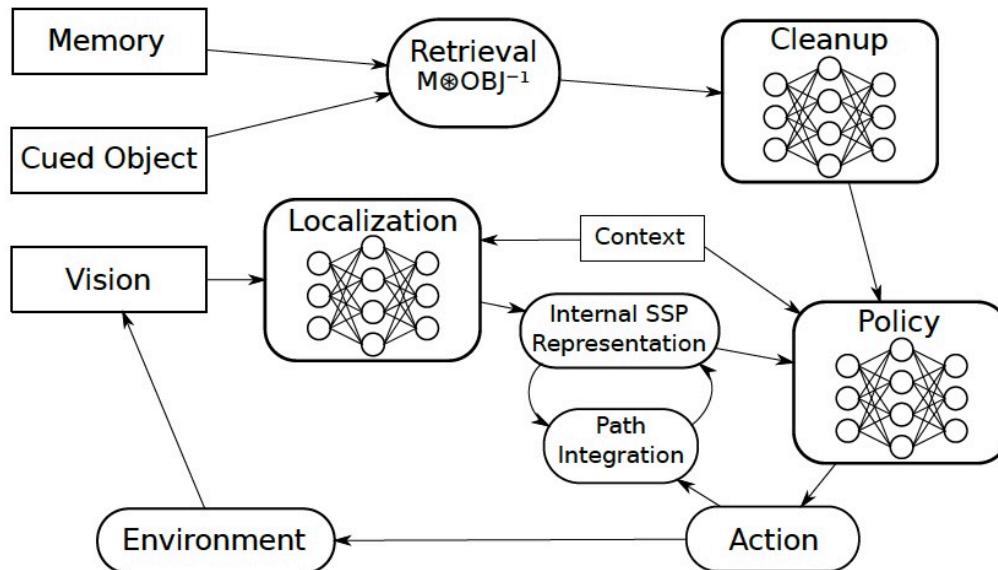


DIAMOND



Navigation network

Combined all of the above into a network to recall location and navigate to arbitrary objects in a maze.



SSPs as Probabilities

- SSPs can be a method for encoding and processing probabilities
- Method directly connects neural networks to probabilistic reasoning
- SSP based methods are efficient

SSPs as Probabilities

Background

- Kernel Density Estimators

From a dataset

$$\mathcal{D} = (x_1, x_2, \dots, x_n)$$

With a kernel function

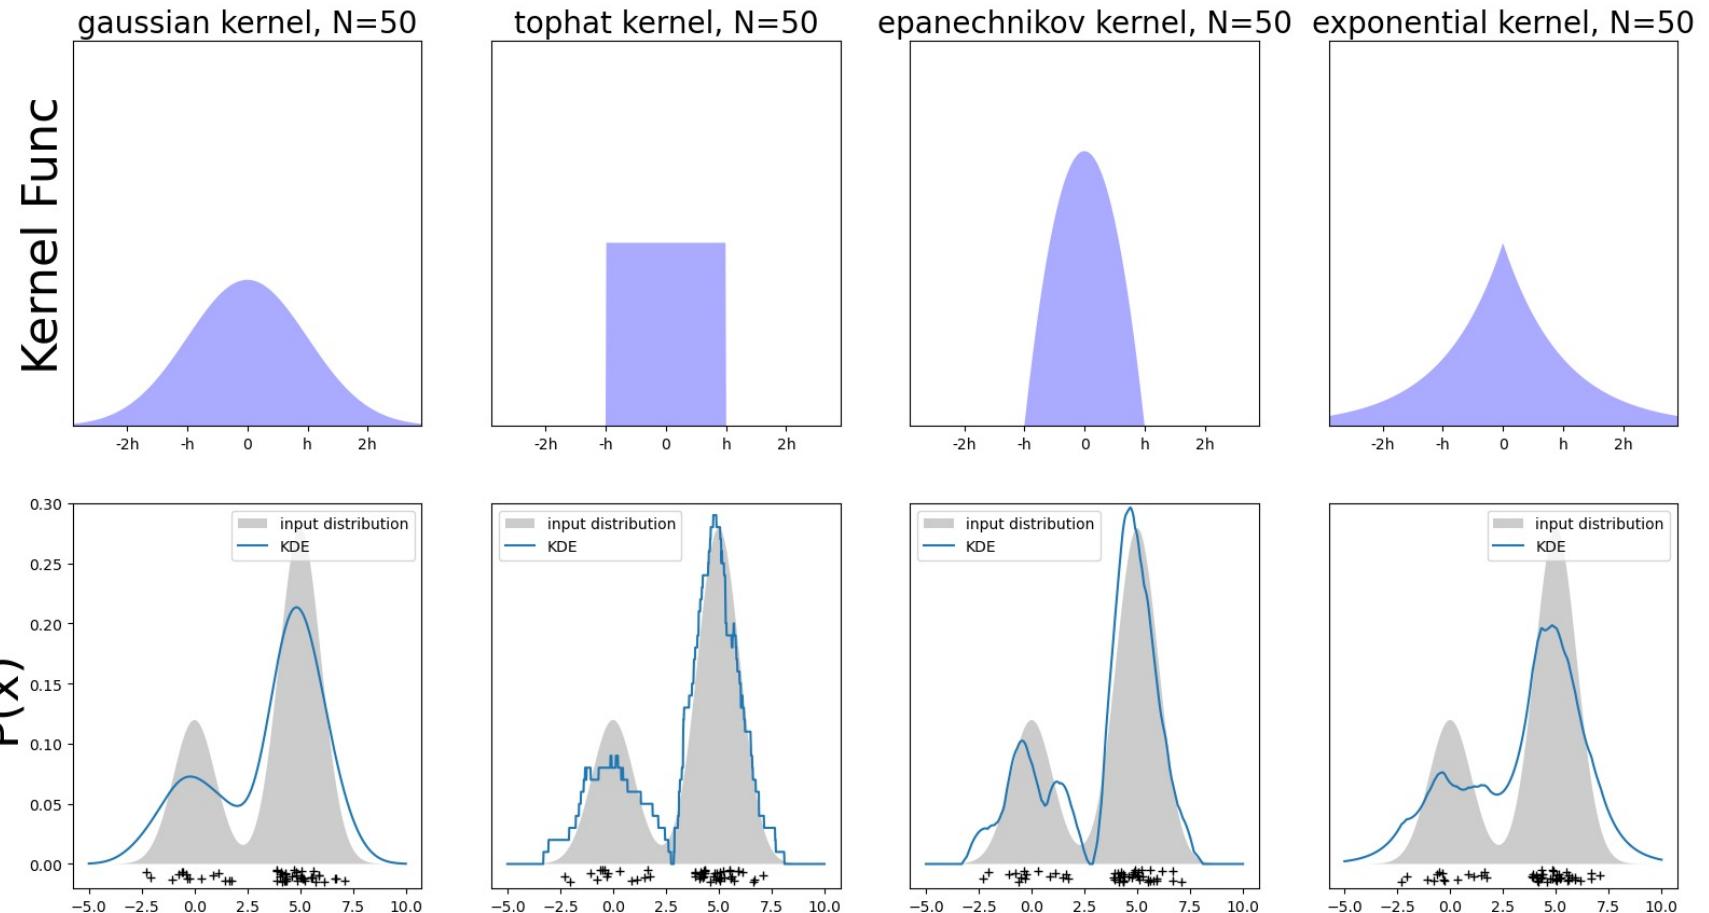
$$k_h(x, x') = k\left(\frac{\|x - x'\|}{h}\right)$$

We can estimate
the probability of x

$$P_{\mathcal{D}}(X = x) = \frac{1}{nh} \sum_{x_i \in \mathcal{D}} k_h(x, x_i)$$

SSPs as Probabilities

- Kernel Density Estimators Examples



SSPs as Probabilities

Problems

- KDE memory grows linearly with the number of observations
- KDE time to compute a probability grow linearly with # of observations

But...

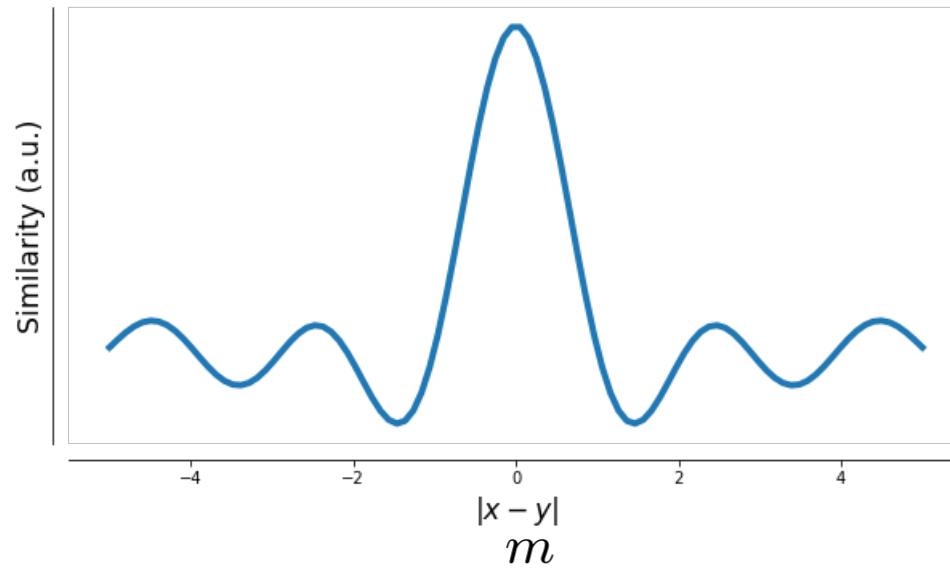
- Not if your kernel is a dot product

$$k_h(\mathbf{x}, \mathbf{x}') \approx \phi_h(\mathbf{x}) \cdot \phi_h(\mathbf{x}') \implies P(\mathbf{X} = \mathbf{x}) = \frac{1}{nh} \sum_{\mathbf{x}_i \in \mathcal{D}} \phi_h(\mathbf{x}) \cdot \phi_h(\mathbf{x}_i)$$

SSPs as Probabilities

SSPs induce a quasi-kernel

- ‘quasi’ because there are negatives



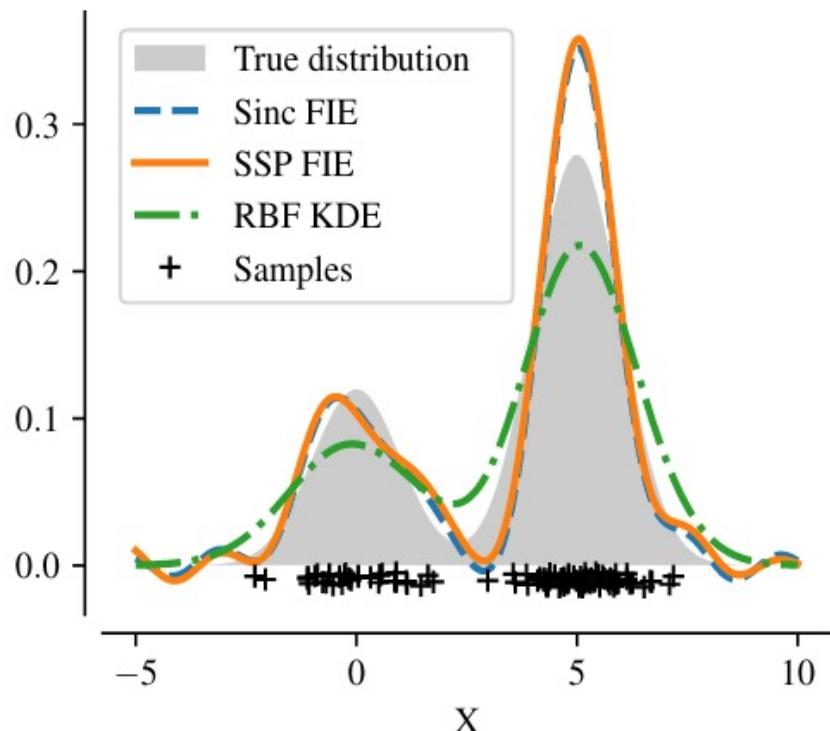
$$\phi(\mathbf{x}) \cdot \phi(\mathbf{y}) \approx \prod_{k=1} \text{sinc}(|x_k - y_k|)$$

SSPs as Probabilities

Can convert to probability estimator

$$P_{\mathcal{D}}(X = x) = \max \{0, \phi(x) \cdot M_{\mathcal{D}, h} - \xi\}$$

Glad et al, 2003



$$\text{ReLU}(\mathbf{w} \cdot \mathbf{z} + b)$$

SSPs as Probabilities

So SSP memory is a latent probability distribution

- The distribution is stored in bundles of vector symbols.
- We can apply manipulations to bundles to produce probabilistic statements

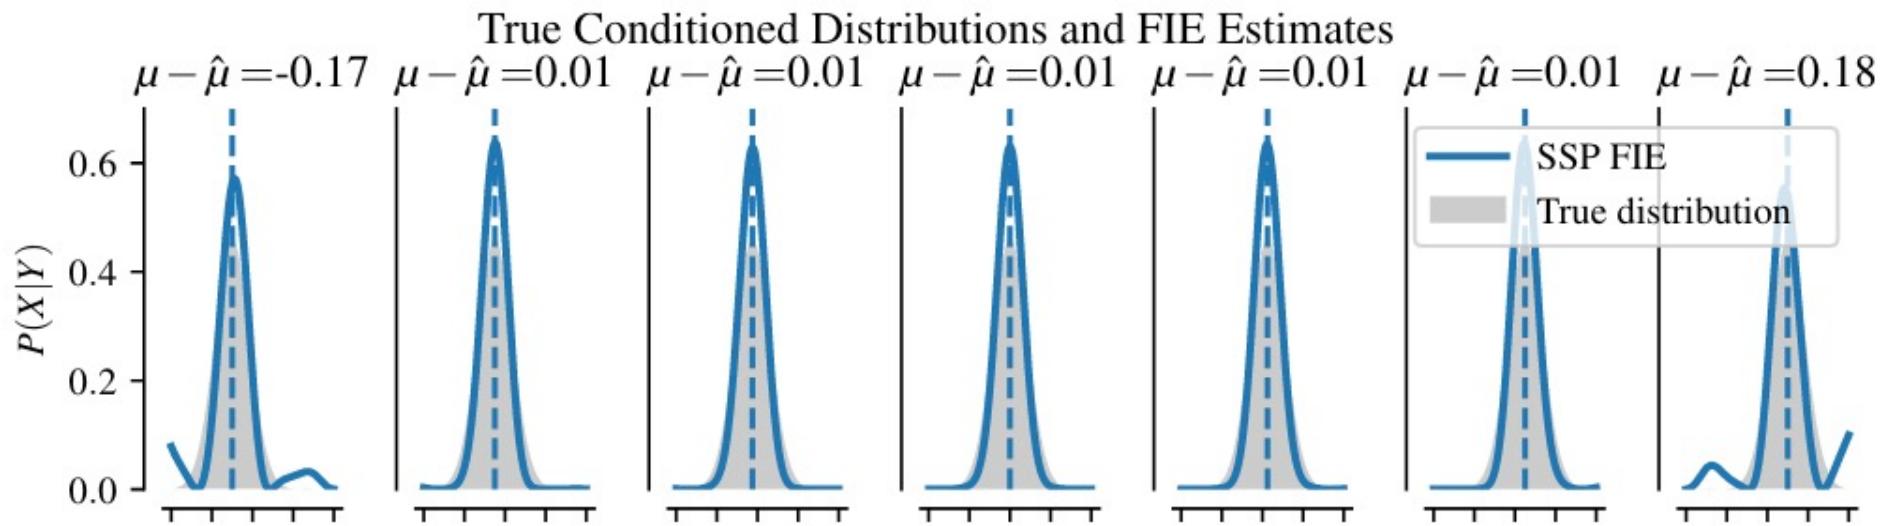
$$M_{\mathcal{D}} = \frac{1}{n} \sum_{\mathbf{x}_i \in \mathcal{D}} \phi(\mathbf{x}_i)$$

$$P(\mathbf{X} = \mathbf{x}) = \phi_h(\mathbf{x}) \cdot M_{\mathcal{D}}$$

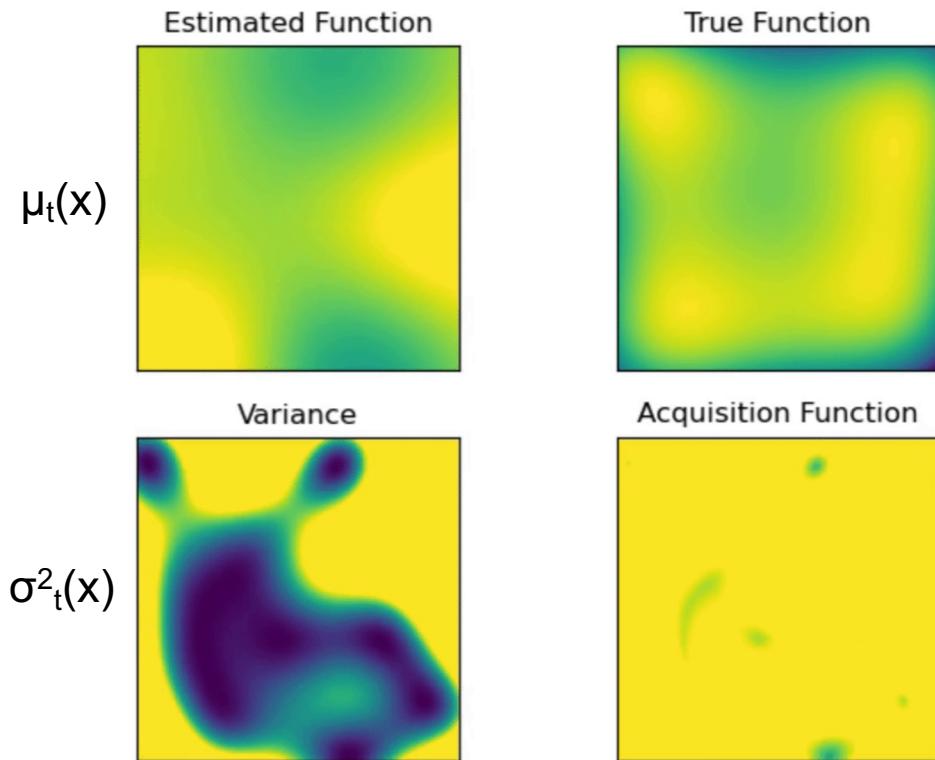
SSPs as Probabilities

Conditioning

$$P(X = x \mid Y = y) \approx \phi_X(x) \cdot [M_{\mathcal{D}} \circledast \phi^{-1}(y)]$$



SSPs as Probabilities

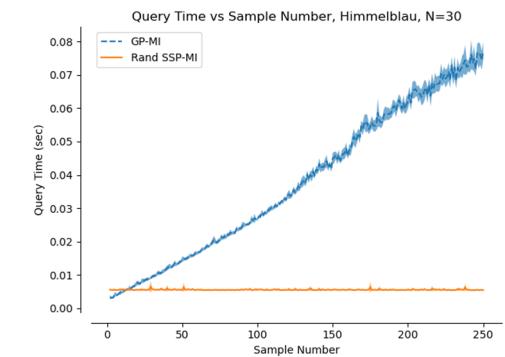
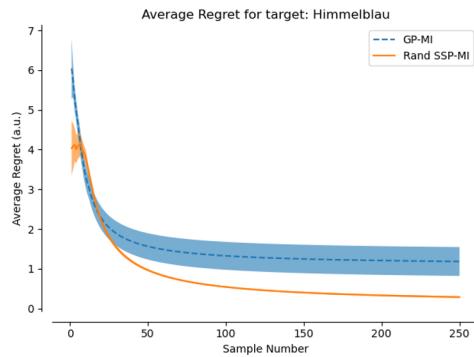
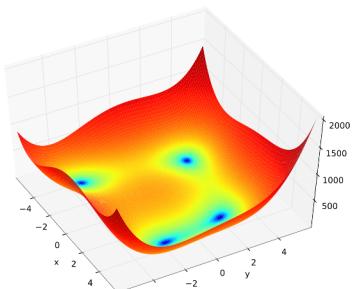


- Mutual Information (MI) is a common objective function used in exploration
- Gaussian Processes (GPs) are a convenient, but computationally intensive tool for computing MI
- We use Spatial Semantic Pointers and Bayesian linear regression to approximate a GP while improving in memory and time complexity

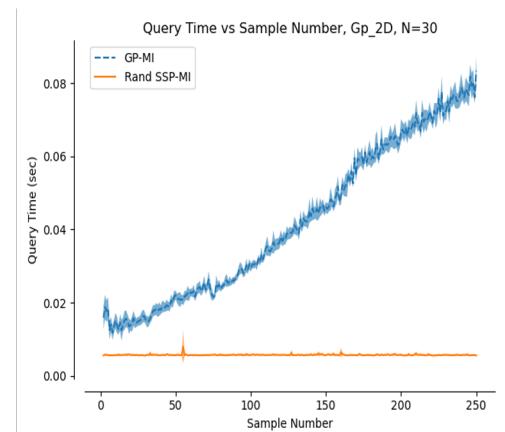
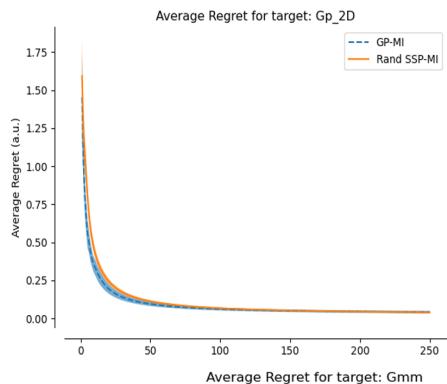
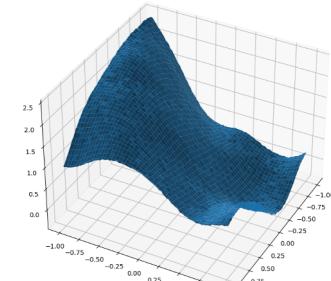
$$MI(X; Y) = H(Y) - H(Y|X) \text{ where } H(Y) = -\sum p(y) \log(p(y))$$

Results

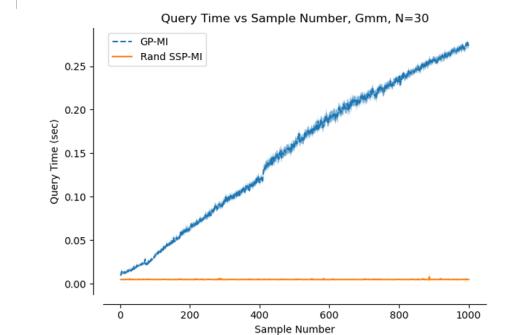
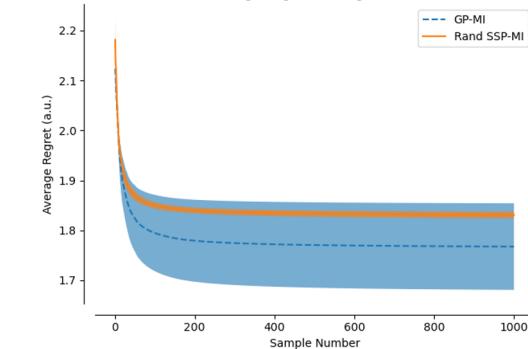
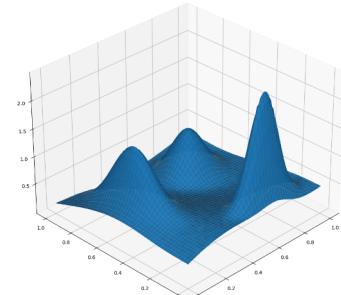
Himmelblau
Function



Sampled from
GP with Matern
Kernel + 1%
noise



GMM with noise
~ (Matern Kernel
+ 1% noise)



SSPs as Probabilities

Where we differ:

- Provide a general and abstract framework for modelling probabilities
- Draw a direct connection between cognitive models and probability statements
- Provide network architectures for conditioning, marginalization, entropy, and mutual information

Conclusion

- SSPs support a variety of types of inference for cognitive models
 - Binding spatial and ‘symbolic’ representations
 - Representations of sampled data that can be used for probabilistic inference
- Improves
 - Interpretability
 - Efficiency (SWaP critical)
- Best engineered systems continue to learn from understanding biological solutions